

Question 1: Entity Relationship Diagram (ERD)

The goal is to draw a Logical ERD using UML notation, which must include Primary Keys (PKs), Foreign Keys (FKs), and resolve the many-to-many (M:N) relationship¹. All entities must have surrogate PKs².

Steps for the Perfect ERD:

1. Identify Entities and Attributes:

- **Category:** CategoryID (PK, surrogate), Description³³³³.
- **Medication:** MedicationID (PK, surrogate), Name⁴.
- **ActiveIngredient:** ActiveIngredientID (PK, surrogate), Name⁵.

2. Resolve the M:N Relationship (Medication and Active Ingredient):

- The relationship is M:N ("one or more active ingredients," "more than one medication") and requires storing a descriptive attribute (Quantity)⁶.
- Create an **Associative Entity: MedicationActiveIngredient**.
- Its Primary Key must be a **Composite Key** made of the two Foreign Keys: MedicationID (PK, FK) and ActiveIngredientID (PK, FK).
- Add the descriptive attribute: Quantity.

3. Establish 1:M Relationships and Foreign Keys:

- **Medication and Category:** Many Medications belong to one Category⁷. Add **CategoryID (FK)** to the Medication entity.
- **ActiveIngredient and Category:** Many Active Ingredients belong to one Category⁸. Add **CategoryID (FK)** to the ActiveIngredient entity.

4. Determine Multiplicities (Cardinality):

- **Category (1) \$\leftrightarrow\$ Medication (1..*):** One Category can have one or more Medications⁹.
- **Category (1) \$\leftrightarrow\$ ActiveIngredient (1..*):** One Category can have one

or more Active Ingredients.

- **Medication (1)** \rightarrow **MedicationActiveIngredient (1..*)**: One Medication can contain *one or more* records in the associative table¹⁰.
 - **ActiveIngredient (1)** \rightarrow **MedicationActiveIngredient (1..*)**: One Active Ingredient can be used in *one or more* records in the associative table¹¹.
-

Question 2: Normalisation to 2NF and 3NF

This question requires normalizing a 1NF table extract to 2NF (Q.2.1) and then to 3NF (Q.2.2), showing all steps, explanations, and final answers in dependency diagram format¹²¹²¹².

Initial 1NF Table and Dependencies

To achieve the required 2NF step, we must assume a **Composite Key**. Based on the data's logical purpose, the assumed Composite Primary Key is **(PharmacyID, PharmacistID)**¹³.

Attribute	Dependency	Type of Dependency
Composite Key	(PharmacyID, PharmacistID)	
Non-Key Attributes	PharmacyName, GroupName, GroupID, PharmacistName	
Partial Dependencies (Violation of 2NF)	PharmacyID \rightarrow PharmacyName, GroupID, GroupName (Determined by only part of the key)	Partial Dependency

	PharmacistID \$\rightarrow\$ PharmacistName (Determined by only part of the key)	Partial Dependency
Transitive Dependency (Violation of 3NF)	GroupID \$\rightarrow\$ GroupName (A non-key attribute determines another non-key attribute)	Transitive Dependency

Q.2.1: Normalise to Second Normal Form (2NF)

Goal: Resolve **Partial Dependencies**.

Explanation: In 1NF, a non-key attribute is dependent on only *part* of the composite primary key. To reach 2NF, we remove these partial dependencies by creating new tables where the determinant (the part of the key) becomes the new Primary Key¹⁴.

Diagram	Table Name	Relationship	Explanation
Pharmacist (PharmacistID \$\rightarrow\$ PharmacistName)	Pharmacist	PharmacistID (PK) \$\rightarrow\$ PharmacistName	Creates a table for pharmacist details, removing the partial dependency of PharmacistName on PharmacistID.
Pharmacy (PharmacyID \$\rightarrow\$ PharmacyName, GroupID, GroupName)	Pharmacy	PharmacyID (PK) \$\rightarrow\$ PharmacyName, GroupID, GroupName	Creates a table for pharmacy location and group details, removing the partial dependency of these attributes on PharmacyID.

Original Table (Junction)	Pharmacy-Pharmacist	PharmacyID (PK, FK) + PharmacistID (PK, FK)	The original table is reduced to a junction table to link pharmacies and pharmacists.
----------------------------------	----------------------------	---------------------------------------------	---------------------------------------------------------------------------------------

Q.2.2: Normalise to Third Normal Form (3NF)

Goal: Resolve Transitive Dependencies.

Explanation: In the 2NF Pharmacy table, the non-key attribute GroupID determines the non-key attribute GroupName ($\text{GroupID} \rightarrow \text{GroupName}$). This is a transitive dependency, violating 3NF. To resolve it, we create a new table for the transitive relationship and leave a Foreign Key in the original table¹⁵.

Diagram	Table Name	Relationship	Explanation
New Table (Group)	Group	GroupID (PK) \$\rightarrow\$ GroupName	Creates a separate table for the group information, using GroupID as the PK.
Updated Pharmacy Table	Pharmacy	PharmacyID (PK) \$\rightarrow\$ PharmacyName, GroupID (FK)	GroupName is removed, leaving GroupID as a Foreign Key to link to the new Group table.

Final 3NF Dependency Diagrams
Group: GroupID (PK) \$\rightarrow\$ GroupName
Pharmacist: PharmacistID (PK) \$\rightarrow\$ PharmacistName

Pharmacy: PharmacyID (PK) \rightarrow PharmacyName, GroupID (FK)

Pharmacy-Pharmacist: PharmacyID (PK, FK) + PharmacistID (PK, FK)

Question 3: Practical SQL Script (MySQL)

The solution must be a single SQL script using the corrected statements below. Use comments to indicate which part answers which question¹⁶.

SQL

-- Ensure the schema is created and used

```
CREATE SCHEMA IF NOT EXISTS Hospital_db;  
USE Hospital_db;
```

-- Q.3.1: Write an SQL statement to create the Patient table. (5 Marks)

```
CREATE TABLE Patient(  
    PatientID INT AUTO_INCREMENT NOT NULL,  
    PatientName VARCHAR(250) NOT NULL, -- Per ERD extract [cite: 257]  
    PatientSurname VARCHAR(250) NOT NULL, -- Per ERD extract [cite: 257]  
    PatientDOB DATE NOT NULL,  
    PRIMARY KEY (PatientID)  
);
```

-- Q.3.2: Write an SQL statement to create the Doctor table. (4 Marks)

```
CREATE TABLE Doctor(  
    DoctorID INT AUTO_INCREMENT NOT NULL,  
    DoctorName VARCHAR(250) NOT NULL, -- Per ERD extract [cite: 261]  
    DoctorSurname VARCHAR(250) NOT NULL, -- Per ERD extract [cite: 262]  
    PRIMARY KEY (DoctorID)  
);
```

-- Q.3.3: Write an SQL statement to create the Appointments table. (9 Marks)

```
CREATE TABLE Appointments(  
    ...  
);
```

```
AppointmentID INT AUTO_INCREMENT NOT NULL,  
PatientID INT NOT NULL, -- Foreign Key [cite: 266]  
DoctorID INT NOT NULL, -- Foreign Key [cite: 267]  
AppointmentDate DATE NOT NULL,  
AppointmentTime VARCHAR(250) NOT NULL, -- Using VARCHAR as per the submitted file for  
time/duration data [cite: 377]  
AppointmentDuration VARCHAR(250) NOT NULL,  
PRIMARY KEY (AppointmentID),  
FOREIGN KEY (PatientID) REFERENCES Patient(PatientID), -- Links to Patient table  
FOREIGN KEY (DoctorID) REFERENCES Doctor(DoctorID) -- Links to Doctor table  
);
```

-- Q.3.4: Write SQL statements to insert the provided data. (11 Marks)

```
-- Insert Patient data [cite: 281, 282]  
INSERT INTO Patient (PatientID, PatientName, PatientSurname, PatientDOB) VALUES  
(1, 'Debbie', 'Theart', '1980-03-17'),  
(2, 'Thomas', 'Duncan', '1976-08-12');
```

```
-- Insert Doctor data [cite: 286, 287]  
INSERT INTO Doctor (DoctorID, DoctorName, DoctorSurname) VALUES  
(1, 'Zintle', 'Nukani'),  
(2, 'Ravi', 'Maharaj');
```

```
-- Insert Appointments data (Ensure DoctorID and PatientID values are correct based on the data table)  
[cite: 288, 289]  
INSERT INTO Appointments (AppointmentID, AppointmentDate, AppointmentTime,  
AppointmentDuration, DoctorID, PatientID) VALUES  
(1, '2024-01-15', '9:00', '15', 2, 1), -- Dr 2 (Ravi), Pt 1 (Debbie)  
(2, '2024-01-18', '15:00', '30', 2, 2), -- Dr 2 (Ravi), Pt 2 (Thomas)  
(3, '2024-01-20', '10:00', '15', 1, 1), -- Dr 1 (Zintle), Pt 1 (Debbie)  
(4, '2024-01-21', '11:00', '15', 2, 1); -- Dr 2 (Ravi), Pt 1 (Debbie)
```

-- Q.3.5: Write an SQL statement to display all appointments between 2024-01-16 and 2024-01-20
(inclusive). (4 Marks)

```
SELECT *  
FROM Appointments  
WHERE AppointmentDate BETWEEN '2024-01-16' AND '2024-01-20';
```

-- Q.3.6: Write a SQL statement to display patient names, surnames, and total appointments, sorted by
appointment count (descending). (6 Marks)

```
SELECT  
P.PatientName,  
P.PatientSurname,
```

```
COUNT(A.AppointmentID) AS NumberOfAppointments
FROM
Patient P
INNER JOIN
Appointments A ON P.PatientID = A.PatientID
GROUP BY
P.PatientID, P.PatientName, P.PatientSurname -- Group by PK and names
ORDER BY
NumberOfAppointments DESC; -- Sort in descending order [cite: 295]
```

-- Q.3.7: Write an SQL statement to display all appointments: date (desc), time, doctor's name/surname, patient's name/surname. (10 Marks)

```
SELECT
A.AppointmentDate,
A.AppointmentTime,
D.DoctorName,
D.DoctorSurname,
P.PatientName,
P.PatientSurname
FROM
Appointments A
INNER JOIN
Doctor D ON A.DoctorID = D.DoctorID -- Correct join on DoctorID
INNER JOIN
Patient P ON A.PatientID = P.PatientID -- Correct join on PatientID
ORDER BY
A.AppointmentDate DESC; -- Sort in descending order [cite: 298]
```

-- Q.3.8: Create a view for patients with appointments with Doctor ID 2, displaying name/surname, sorted by surname (asc). (6 Marks)

```
CREATE VIEW Patients_with_DoctorID2 AS
SELECT
P.PatientName,
P.PatientSurname
FROM
Patient P
INNER JOIN
Appointments A ON P.PatientID = A.PatientID -- Correct join key
WHERE
A.DoctorID = 2
ORDER BY
P.PatientSurname ASC; -- Sort in ascending order [cite: 301]
```

-- Q.3.9: Create a stored procedure called get_appointments (IN date) to display appointment details

for that date (time asc, duration, doctor name, patient name). (10 Marks)

DELIMITER //

```
CREATE PROCEDURE get_appointments (
    IN app_date DATE
)
BEGIN
    SELECT
        A.AppointmentTime, -- In ascending order [cite: 309]
        A.AppointmentDuration,
        D.DoctorName,
        D.DoctorSurname,
        P.PatientName,
        P.PatientSurname
    FROM
        Appointments A
    INNER JOIN
        Doctor D ON D.DoctorID = A.DoctorID
    INNER JOIN
        Patient P ON P.PatientID = A.PatientID
    WHERE
        A.AppointmentDate = app_date -- Filter by the input parameter
    ORDER BY
        A.AppointmentTime ASC; -- Sort by time ascending [cite: 309]
END //
```

DELIMITER ;

Question 4: MongoDB Shell Commands

The solution must be a single text file containing the interactive shell commands¹⁷.

Correct MongoDB Commands:

Question	Command (Replace <your-student-number>)	Explanation
Q.4.1 (Create Database)	use patients_<your-student-number>	The use command creates the database if it doesn't exist and switches the context to it ¹⁸ .
Q.4.2 (Create/Insert Data)	db.patients.insertMany([{ "Patient Name": "Debbie", "Patient Surname": "Theart", "Date of Birth": "1980-03-17" }, { "Patient Name": "Thomas", "Patient Surname": "Duncan", "Date of Birth": "1976-08-12" }])	Creates the patients collection and inserts the required data using a single insertMany command ¹⁹ .
Q.4.3 (Get all notes)	db.patients.find().pretty()	The find() method without arguments returns all documents/notes in the collection ²⁰ . .pretty() formats the output.
Q.4.4 (Query notes born after 1979-01-12)	db.patients.find({ "Date of Birth": { \$gt: "1979-01-12" } }).pretty()	Uses the \$gt operator (greater than) for string comparison, which works correctly for YYYY-MM-DD dates, to find records born after the specified date ²¹ .